

IN2002 Data Structures and Algorithms

Lecture 8 – Advanced Trees

Aravin Naren
Semester 1, 2018/19

Learning Objectives

- Understand and be able to use the data structures binary search trees, splay trees and B-trees
- Be able to understand, apply and develop algorithms to handle the data structures above.
Including:
 - Access to keys in trees
 - Balancing trees
 - Adjusting trees

Balanced Trees

Remember...Analysis of binary search trees

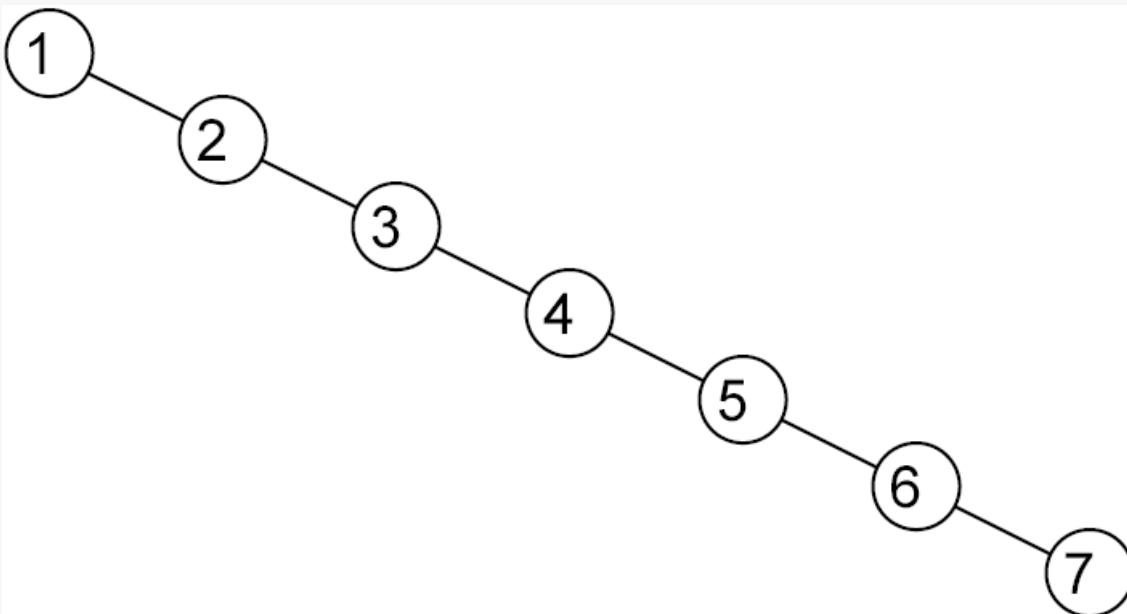
- Search, delete and insert have the height of the tree as the worst case for running time.
- Search, delete and insert take $O(\log n)$, where n is the number of nodes in the tree, if the tree is close to perfectly balanced.
- Otherwise, they can be as bad as $O(n)$
- **Problem:** insertion and deletion make the tree less balanced.
- **Solution:** keep the tree approximately balanced, without making insertion and deletion too expensive

The importance of trees being balanced

- To access a node in a binary search tree, the path is visited from the root to the node.
- This means that the maximum number of nodes visited is equal to the height of the tree.
- In a perfectly balanced tree, this is $\log(n)+1$.

A not-so-balanced tree

- The maximum number of levels of a tree with n node is n .



Should the tree be kept balanced?

The **height h of the tree** will be between $\log(n)+1$ and n , so the time complexity of access to the tree will be between $O(\log n)$ and $O(n)$

Ideally, the tree should be kept **perfectly balanced**, with $h = \log(n) + 1$. Doing this after insertion or deletion takes $O(n)$ time, which defeats the purpose of efficiency.

Solution: keep the tree balanced enough to obtain $O(\log n)$ height.

Some approximate balance definitions

Weight-balanced trees:

each node has $1/k \leq \text{size}(\text{right})/\text{size}(\text{left}) \leq k$

AVL trees (height-balanced):

each node has $-1 \leq \text{height}(\text{right}) - \text{height}(\text{left}) \leq 1$
(AVL-balance)

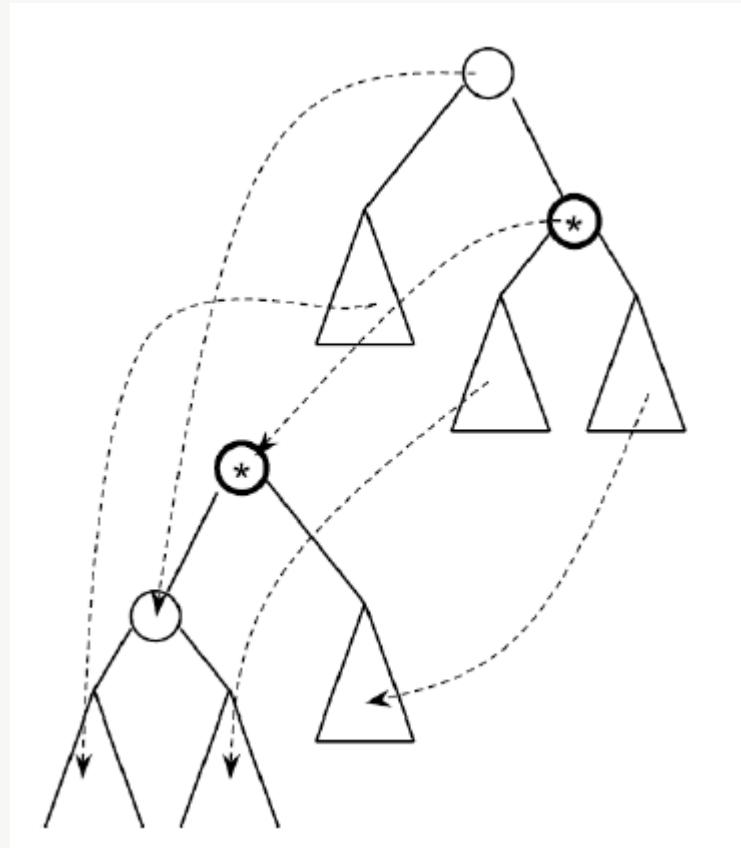
Red-Black trees: each node is either red or black, and no red node can have a red child.

Each full path from the root contains the same number of black nodes.

Balancing Trees

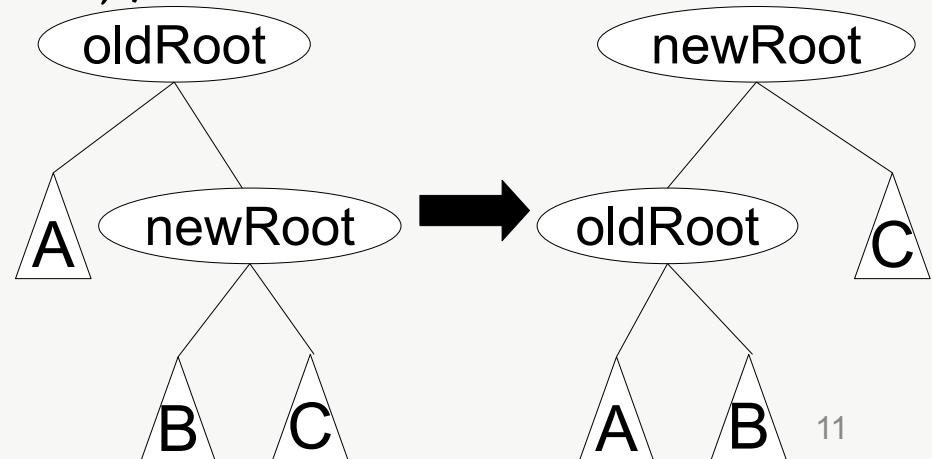
Tree Rotations

- Rotation is a basic operation in many tree re-balancing schemes
- Rotations keep the search tree structure
- Rotations can be applied anywhere in the tree



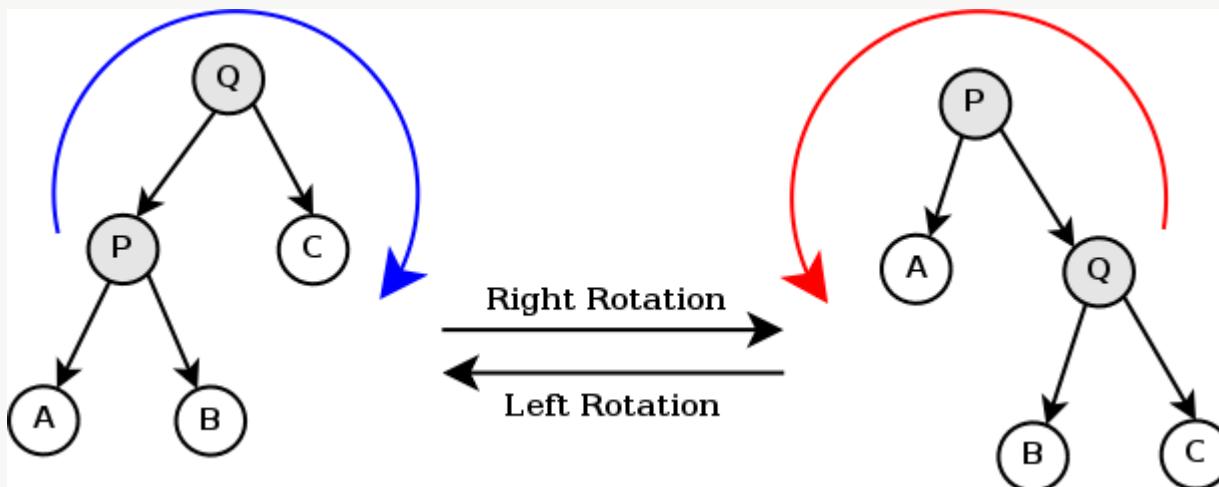
Tree Rotation in Java

```
public TreeNode rotateLeft(TreeNode oldRoot) {  
    TreeNode newRoot = oldRoot.getRight();  
    if(newRoot != null){  
        oldRoot.setRight(newRoot.getLeft());  
        newRoot.setLeft(oldRoot);  
    }  
    return newRoot;  
    // needs to checked  
    // by caller  
}
```

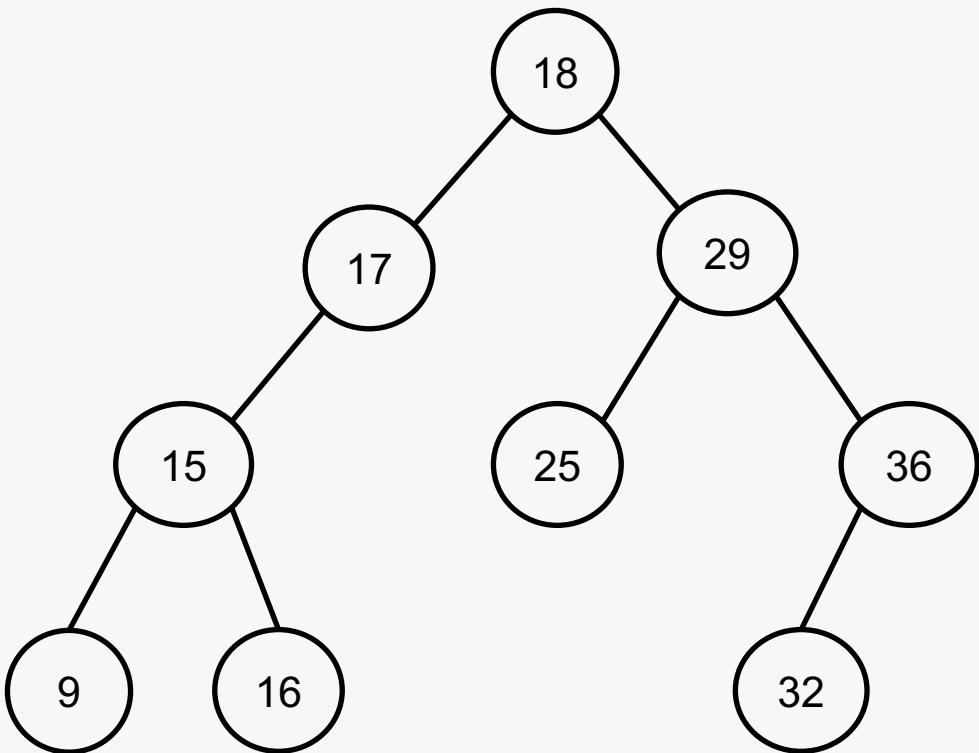


Left and Right Tree Rotations

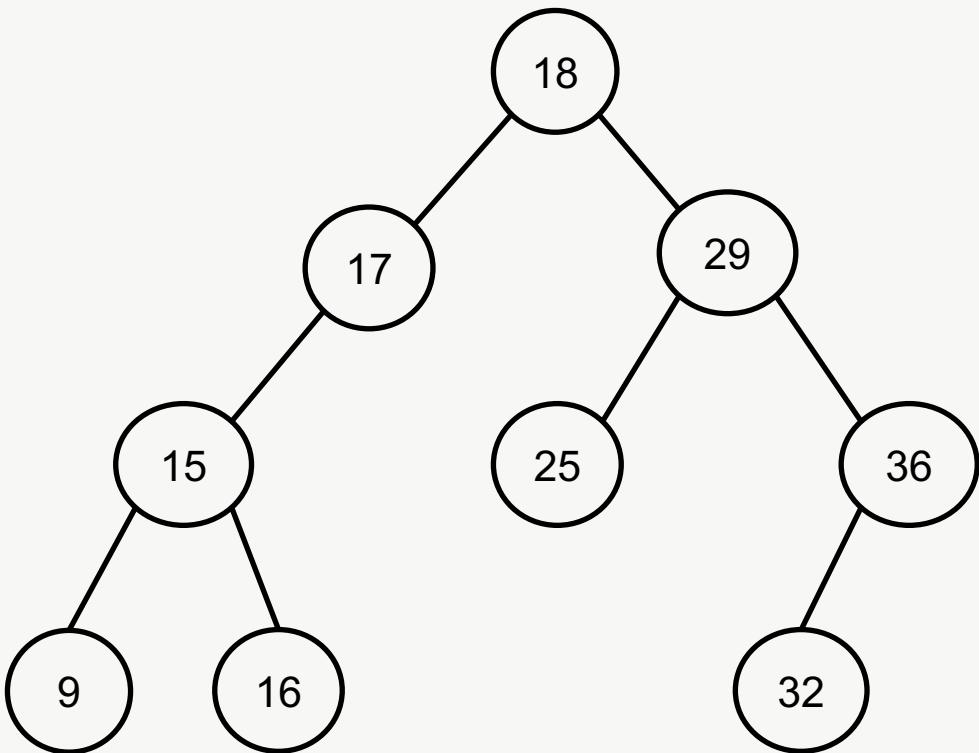
- Rotations can be applied in two directions



Left Tree Rotation Example



Right Tree Rotation Example



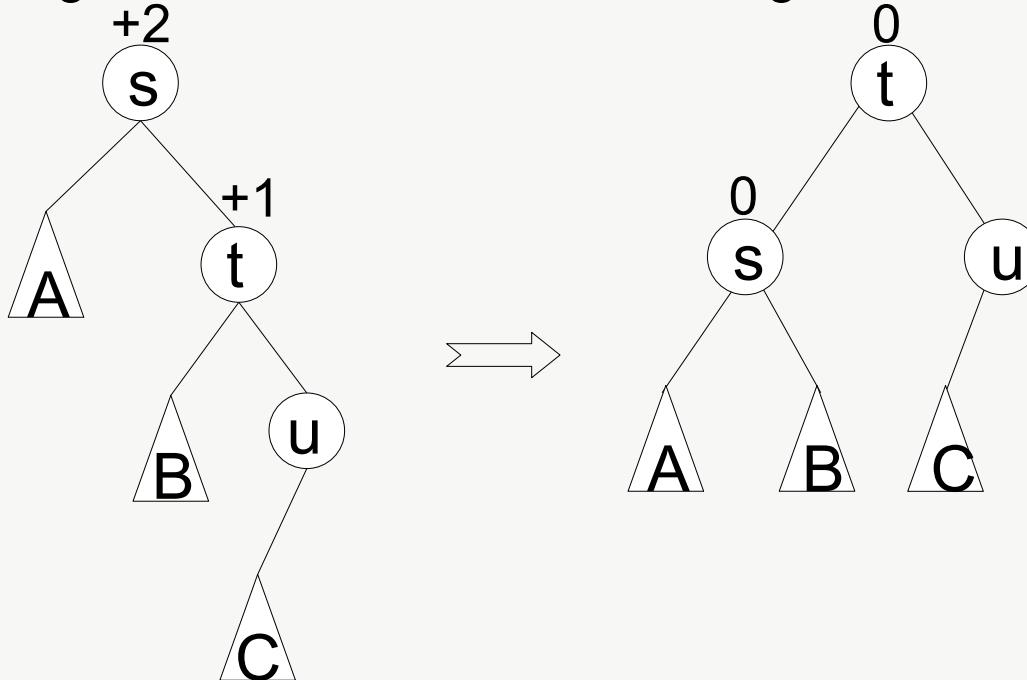
Re-balancing an AVL tree

If insertion invalidates the AVL condition ($-1 \leq \text{height}(\text{right}) - \text{height}(\text{left}) \leq 1$), use rotations to restore it:

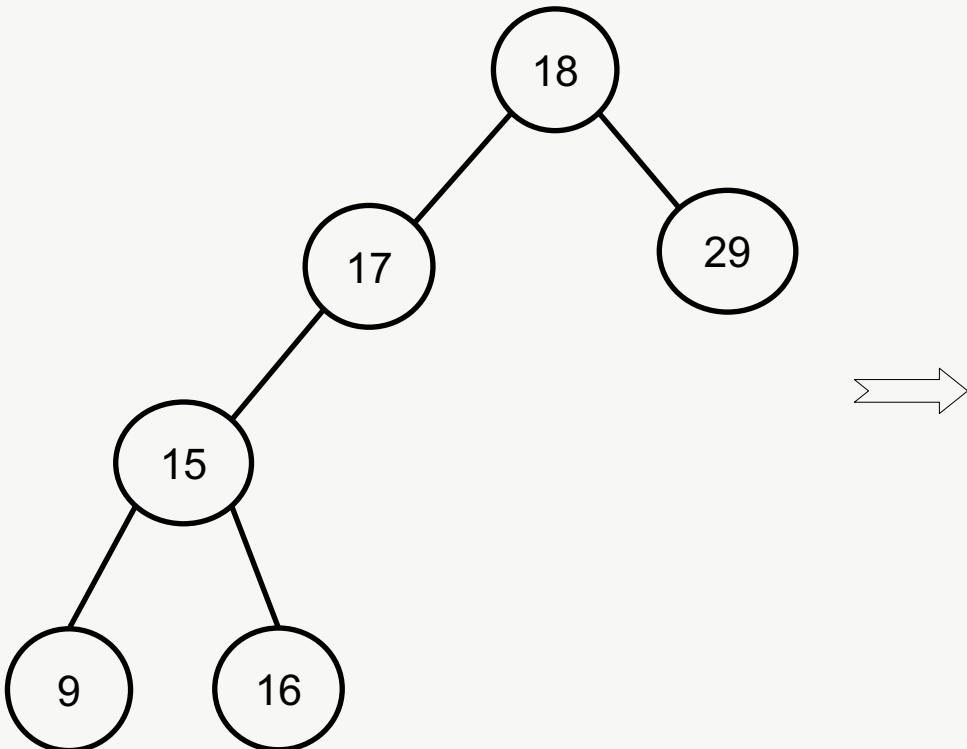
- For insertion, at most two $O(1)$ rotations are required (but finding the position still takes $O(\log n)$)
- Deletions may require rotating any nodes in the path to the changed node, $O(\log n)$
- No re-balancing at lookup-only access

Rebalancing by Single Rotation

If root has AVL-balance < -1 or AVL-balance > 1 and
child balance on higher subtree side has same sign or is 0

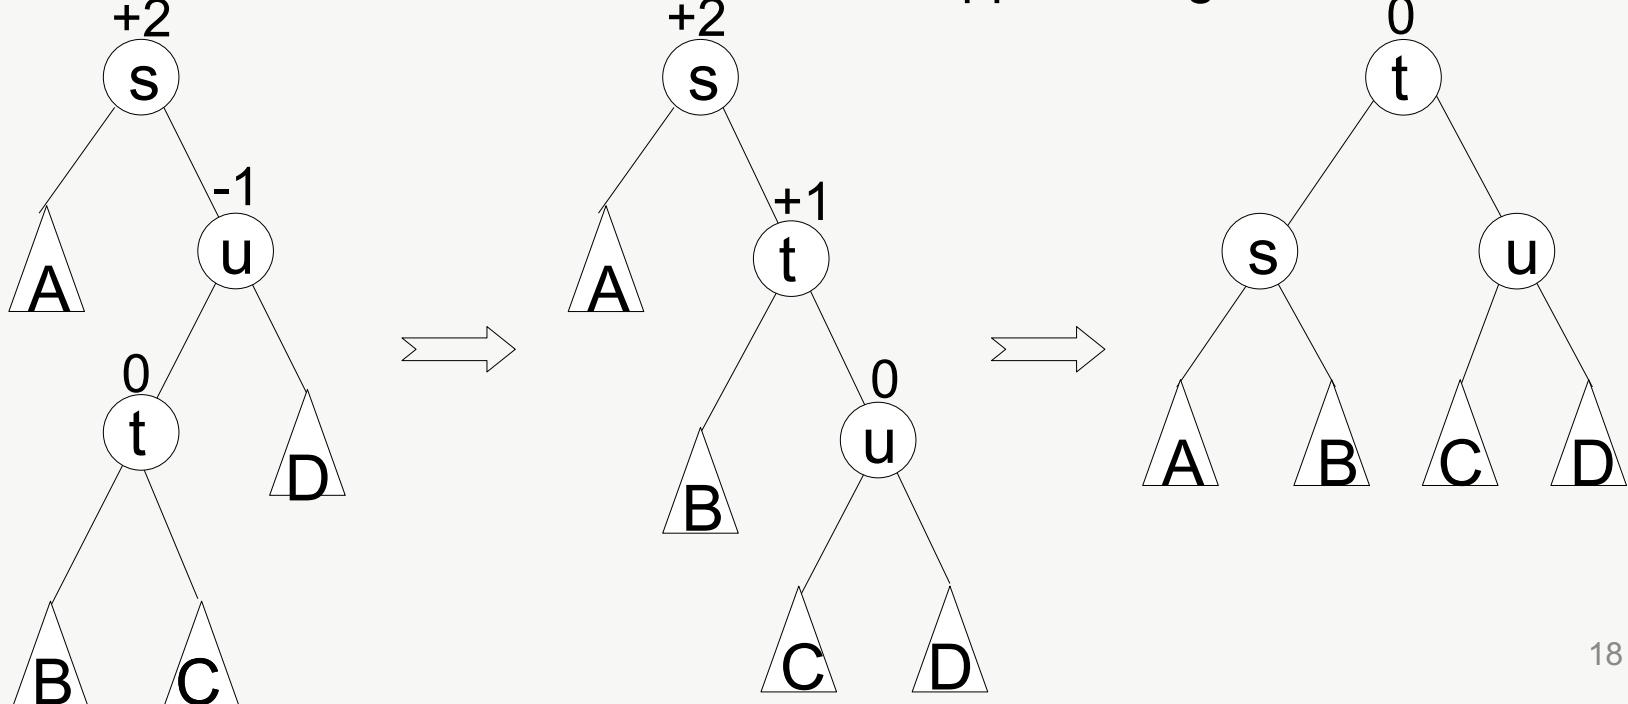


Rebalancing by Single Rotation Example

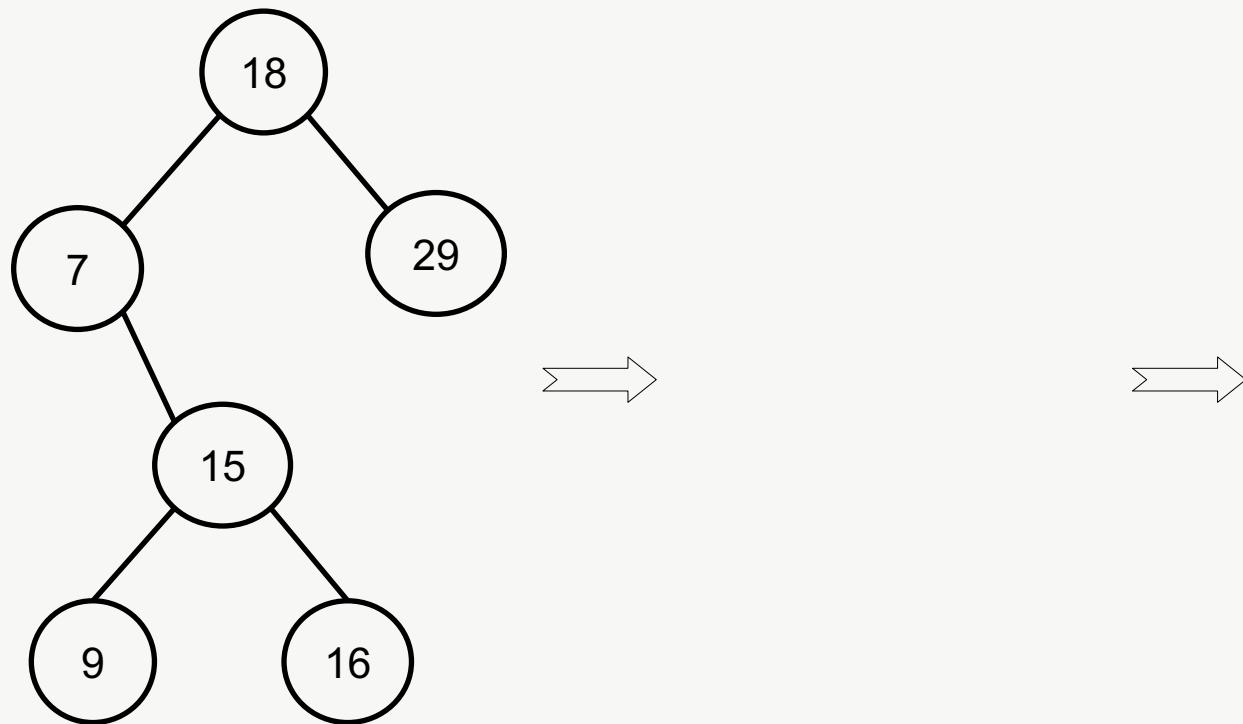


Rebalancing by Double Rotation

If root has AVL-balance < -1 or > 1 and
higher subtree side child AVL-balance has opposite sign



Rebalancing by Double Rotation Example



AVL Tree Rebalancing Animation

- Inserting 12, 23, 200, 17, 15, 20, 18
- Then deleting 15, etc.

AVL Tree

Splay Trees

Splay Trees

- A splay tree rearranges itself after each access
- Moves the accessed node to root (splay operation)
- Frequently accessed keys get closer to the root
- No need for additional information to be stored

Splay step operations

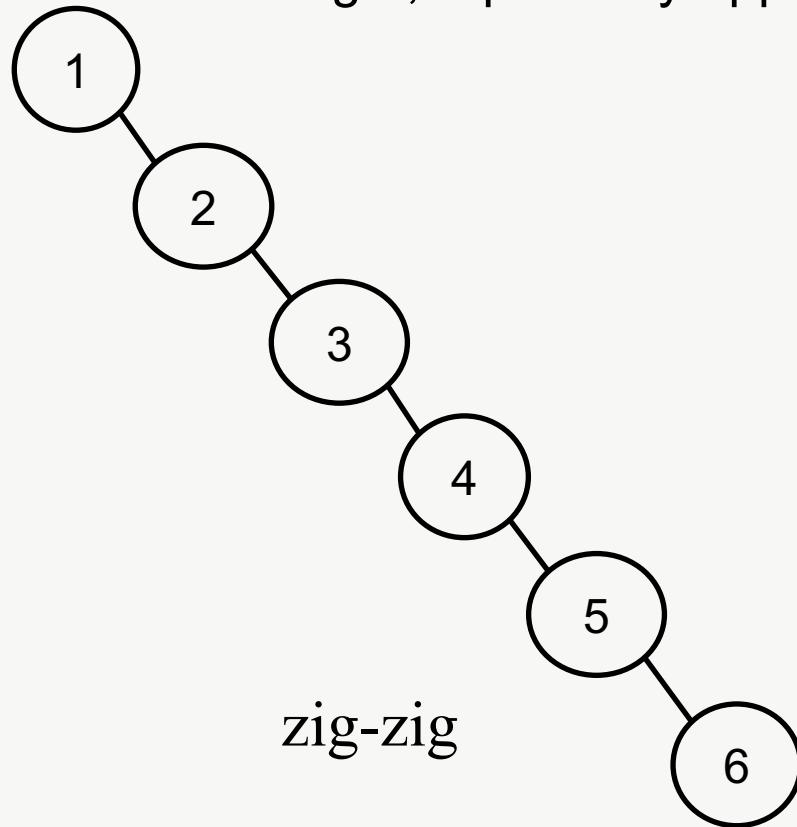
There are three cases (plus a symmetric one), and maximally two rotations per splay:

- Until accessed node is the root:
- If the accessed node's parent is the root (zig), rotate the accessed node up.
- If the accessed node is a left child and its parent is a left child as well (zig-zig), rotate the parent up and then rotate the accessed node up.*
- Otherwise (zig-zag), rotate to the accessed node up twice (standard double rotation).

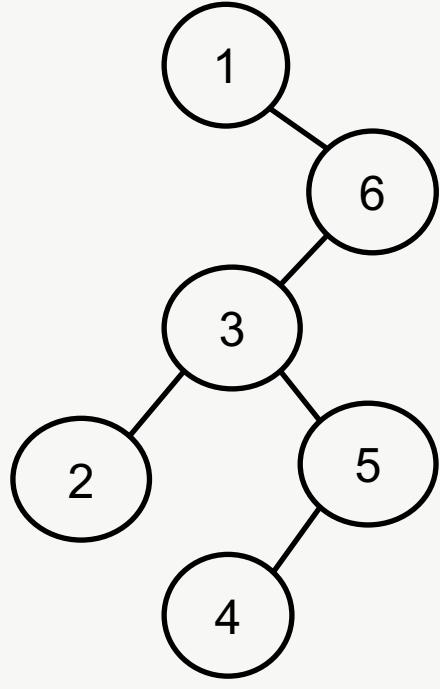
* Also applies to right child and right child

An example splay operation

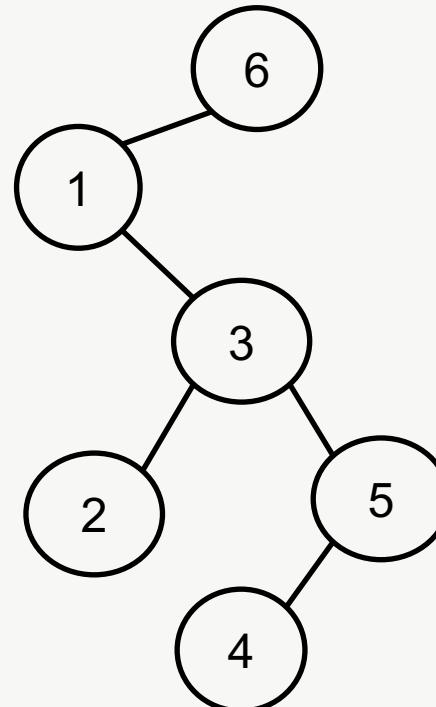
After accessing 6, repeatedly apply splay(6):



Example splay (cntd.)



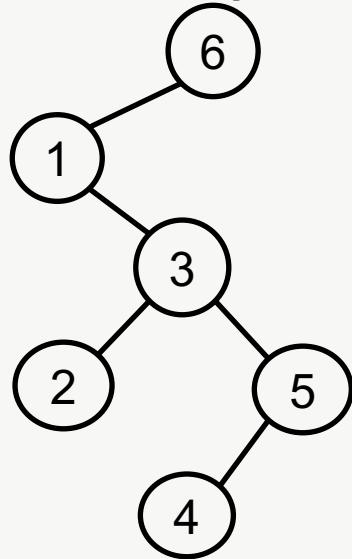
zig



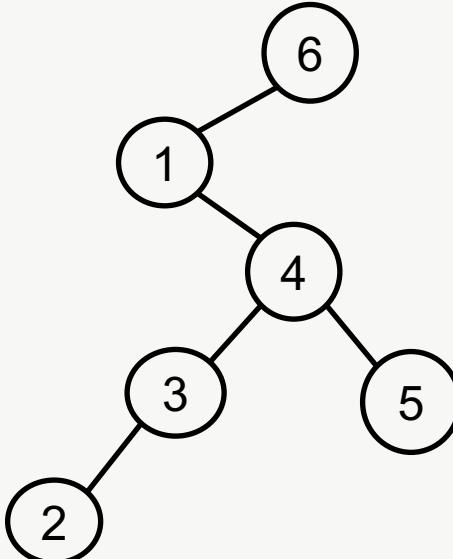
After the third splay, 6 is at the root of the tree.

Example splay zig zag

After accessing 4, repeatedly apply `splay(4)`:

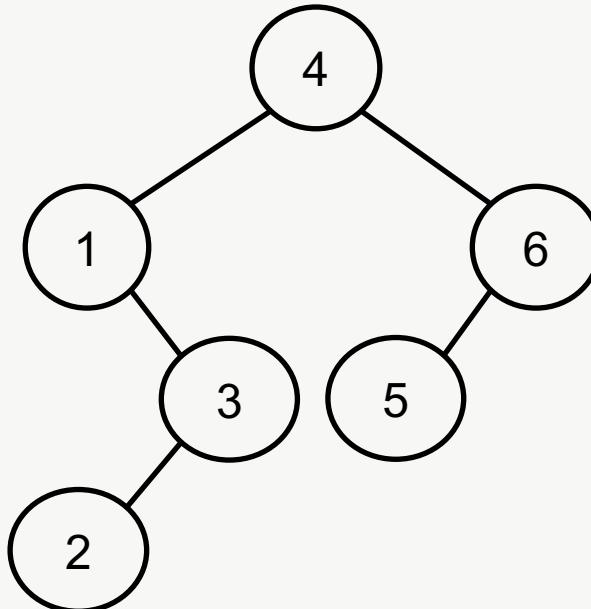


zig-zag



zig-zag

Example splay zig zag (cntd.)



After the second splay, 4 is at the root of the tree.
Note that 6 is still closer to the root than initially.

Splay trees are self-organising

There are three cases (plus a symmetric one), and maximally two rotations per splay:

- Until accessed node is the root:
- If the accessed node's parent is the root (zig), rotate the accessed node up.
- If the accessed node is a left child and its parent is a left child as well (zig-zig), rotate the parent up and then rotate the accessed node up.*
- Otherwise (zig-zag), rotate to the accessed node up twice (standard double rotation).

* Also applies to right child and right child

Splay Tree Animation

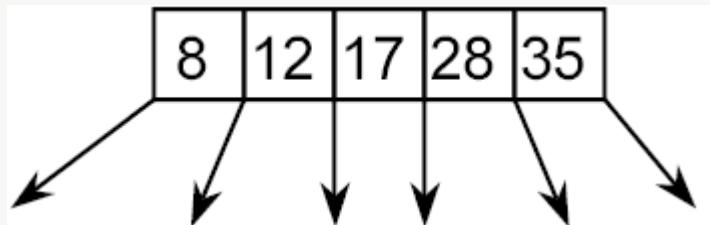
- Inserting 1, 2, 3, 4, 5, 6
- Then finding 1, 4, etc.
- Then deleting ...

Splay tree

Multi-Way Search Trees

Multi-way search trees

Each node has m subtrees, interleaved with $m-1$ keys stored in ascending order



All keys in subtrees to the left of key k are smaller than k

All keys in subtrees to the right of key k are greater

Large m is useful for external storage (hard disk)

Small m is better for internal storage (RAM)

... and binary search trees are multi-way trees with $m = 2$

B-Trees

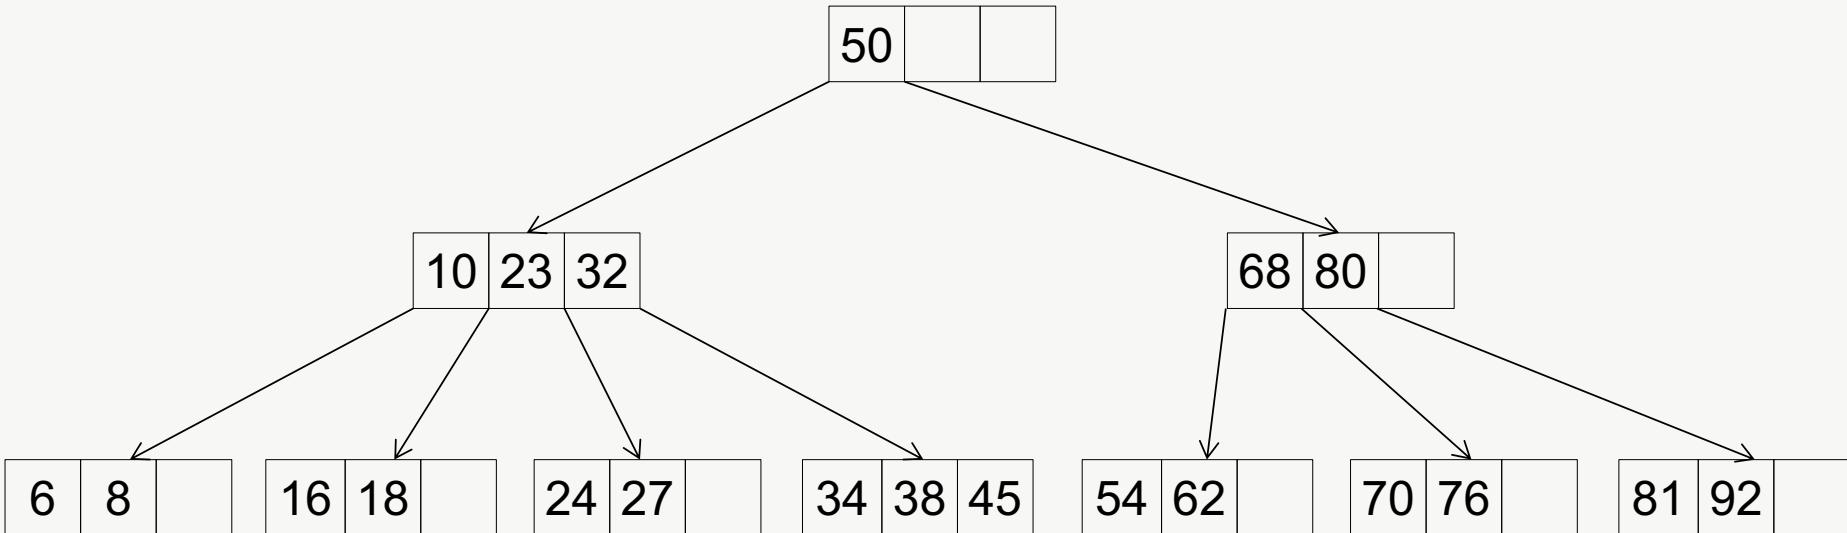
A B-tree of order m is a multi-way search tree for which:

- the root has at least one key
- every non-root node has k subtrees and $k-1$ keys with $m/2 + 1 \leq k \leq m$
- leaves have all subtrees empty
- non-leaf nodes have non-empty subtrees
- all leaves are at the same level

The last condition guarantees that B-trees are approximately balanced.

For hard disk storage, one node is one HD block.

A B-tree of order 4 and height 3



Insertion into a B-tree

Always insert new keys into a leaf

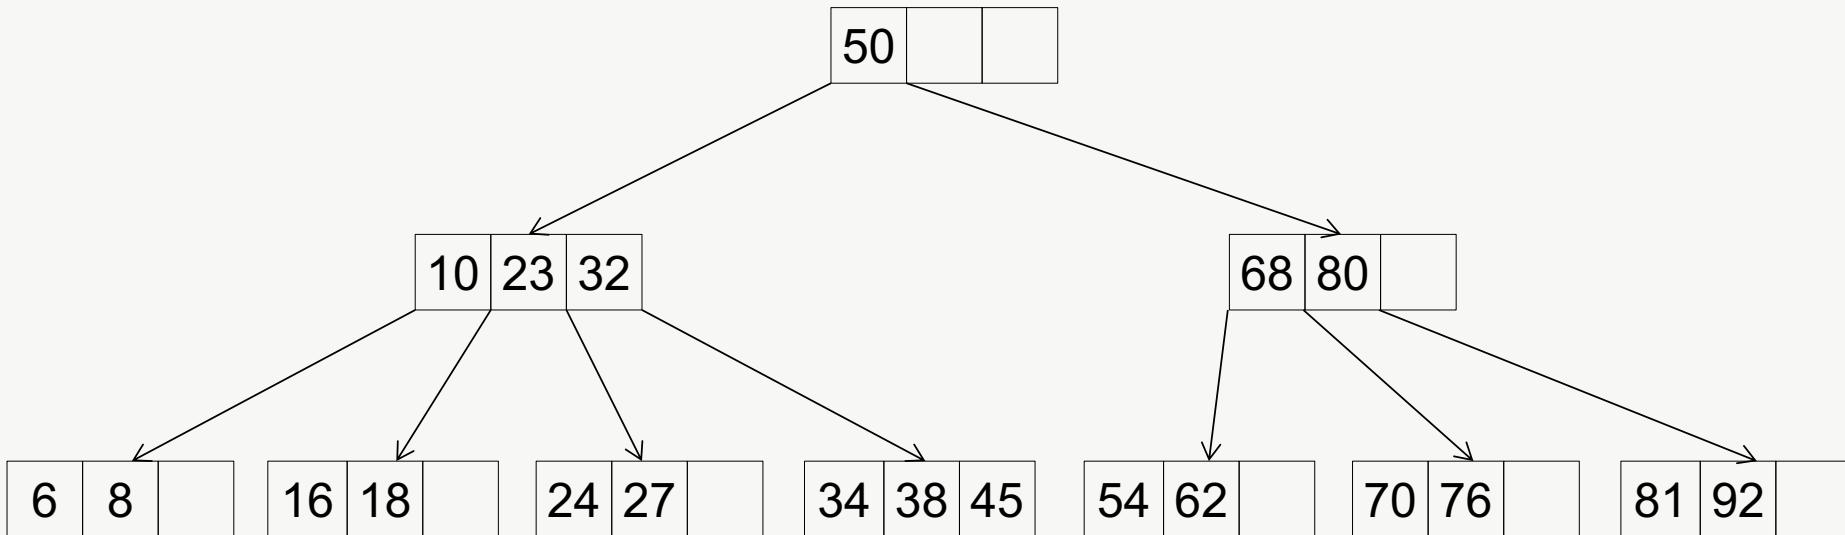
If the node overflows, split it:



- where k is the middle key of the overflowed node
- if the parent node overflows, repeat
- if the root overflows, create a new root on top and split

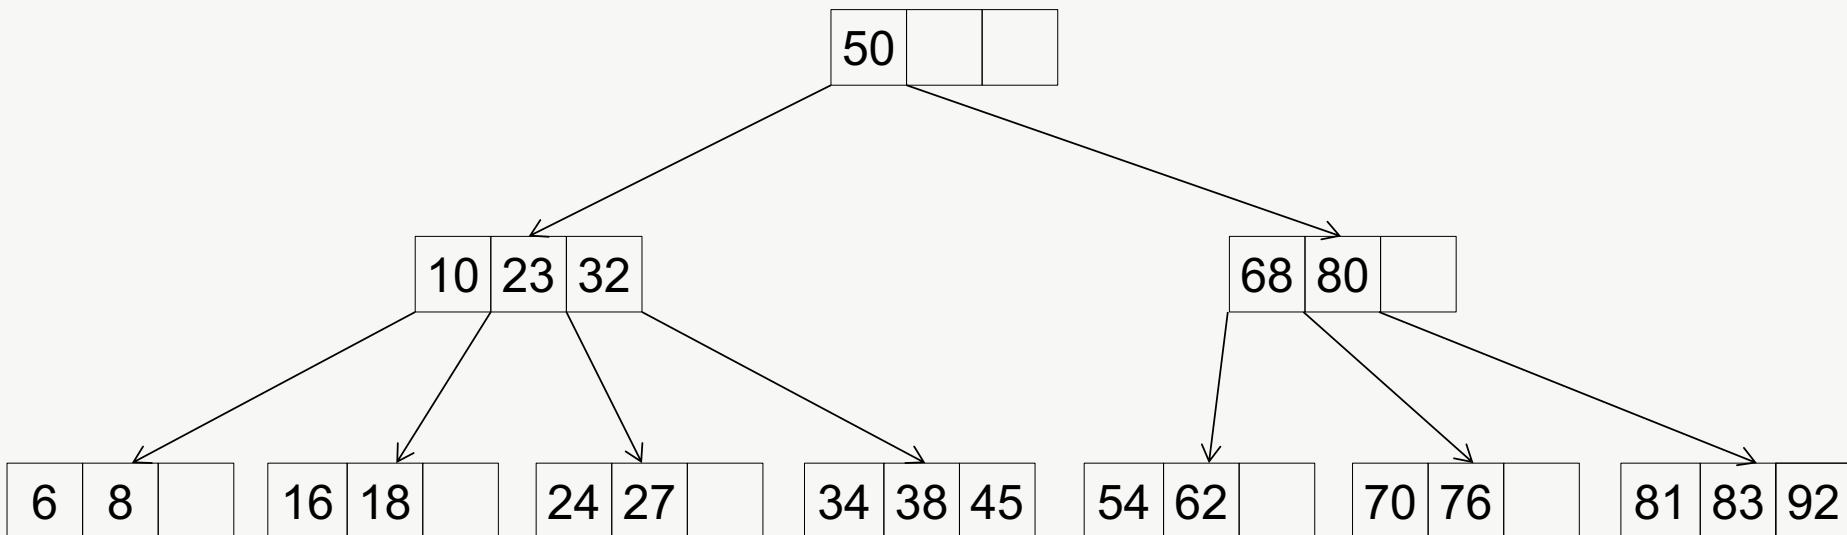
Insertion into a B-tree Example

Insert 83 then 90



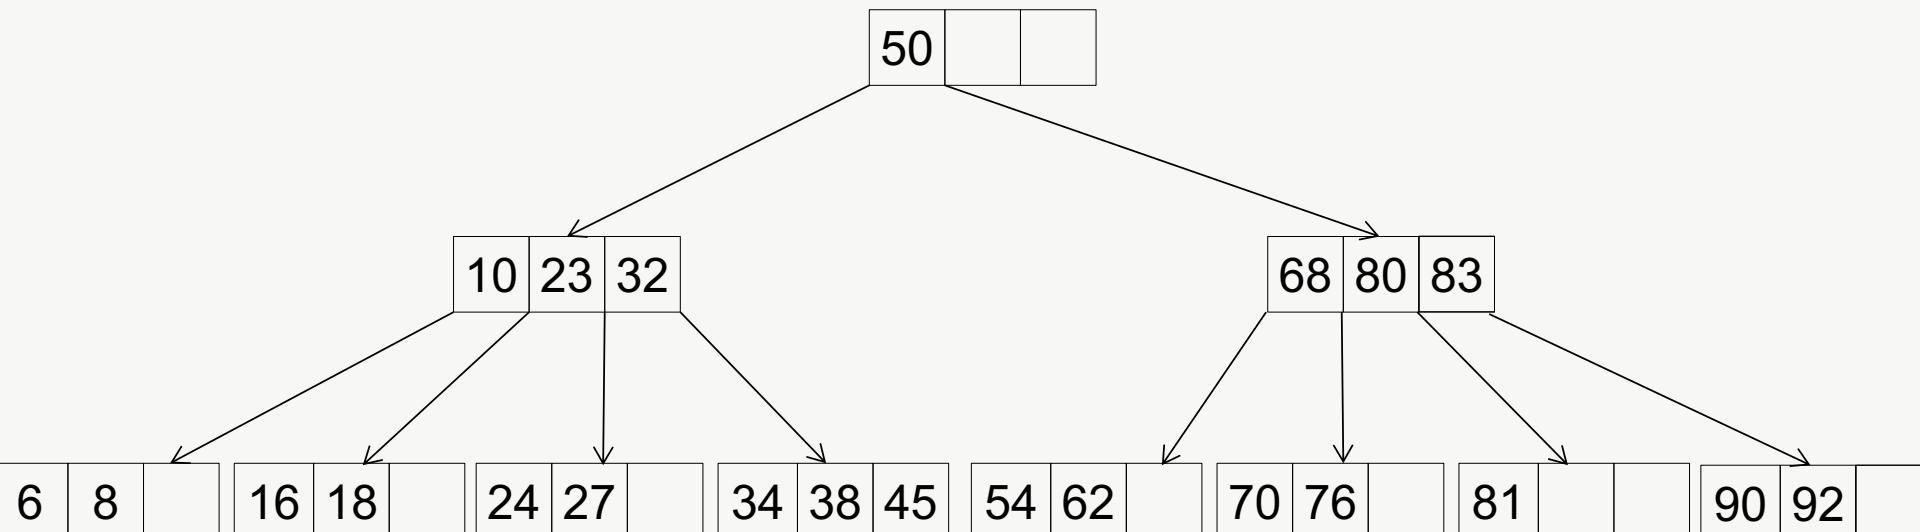
Insertion into a B-tree Example (cntd.)

Insert 83 (then 90)



Insertion into a B-tree Example (cntd.)

Insert 90



Deletion from a B-tree

- If the key to be deleted is in a non-leaf node, swap it with a neighbouring key in a leaf and delete it
- If deletion leaves a node half empty, redistribute keys with a sibling node
- If sibling node is only half full, combine with it (reverse of splitting)
- If the root node ends up with no keys and one pointer, its child becomes the new root

B-Tree Animation

- Inserting ...
- Then finding ...
- Then deleting ...

B-Tree

Reading

- Weiss: Sections 18.4, 18.8, 21.1-21.3, 21.5-21.7
- Drozdek: Sections 6.7, 6.8, and 7.1.1

Next session: Graphs

Drozdek: Sections 8.1, 8.2 and 8.3 OR Weiss: Sections 13.1-13.4

The background image shows a wide-angle aerial view of the London skyline during the day. Key landmarks visible include the London Eye, the River Thames, the Millennium Bridge, the Tate Modern art museum, and various other skyscrapers and historical buildings scattered across the city. The sky is clear and blue.

City, University of London
Northampton Square
London
EC1V 0HB
United Kingdom

T: +44 (0)20 7040 8406

E: ug-compsci@city.ac.uk

www.city.ac.uk/department-computer-science